# edge_inference_onnx_runtime (Initial Draft)

Domain: Edge Iot | Document ID: edge_inference_onnx_runtime_v1
Specification Standard: IEEE Capstone Evaluation Corpus

## Executive Overview:

Optimizing neural inference latency on ARM NEON SIMD instructions using INT8 ONNX Runtime execution providers.

## Methodological Verification & System Lineage:

This specification is catalogued in IntentCloud for intent-aware hybrid retrieval.
Text chunks are indexed into 384-dimensional dense semantic vectors and 1,000,003-dimensional sparse feature-hashed token projections.
Candidate pools are unified using scale-invariant Reciprocal Rank Fusion (k=60).
Explainable citation attribution is extracted using cross-encoder cross-attention.